

Assignment 1: High Performance Dense Linear Solver

Christian Wiskott, 01505394.

April 19, 2020

Part 2: Unblocked Right-looking LU Factorization

In this part of the assignment, the goal was to write an efficient algorithm for calculating an unblocked right-looking LU-Factorization *without pivoting* for a randomly generated, non-singular matrix $A \in \mathbb{R}^{n \times n}$. With this factorization, A is decomposed into an upper triangular matrix U and a lower triangular matrix L , such that $A = LU$. The main idea is to apply Gauß elimination to A , to eliminate the elements below the main diagonal, which transforms A into the upper triangular matrix U . The elements of the matrix L are the factors computed during the Gauß elimination.

The matrix A was generated using a uniform distribution featuring randomly chosen values within the interval of $[90, 100)$.

Performance Evaluation

In order to evaluate the performance of the written code, the run-time was measured, such that only the actual computation of the resulting matrix U was included in the measurement. This computation required three explicit for-loops which scale with $\mathcal{O}(n^3)$ and can be examined in the provided code.

During the computation only one matrix U was used, instead of creating the two matrices U and L . This was made possible by overwriting the given matrix A and adapting the last for-loop, to not overwrite the coefficients otherwise being in L . This choice was motivated by reducing the necessary memory during the computation, since only one matrix had to be stored, but had no significant effect on the run-time of the algorithm.

Fig. 1 shows the run-time analysis of the LU-factorization for problem sizes within the range of $[100, 10000]$. The data has been fitted with a cubic function,

which shows (with high correlation) a cubic increase w.r.t the matrix dimension, which coincides with the expected $\mathcal{O}(n^3)$ -scaling. Due to this scaling, the required computation time increased very rapidly and reached more than 300 seconds for size $10k \times 10k$. This led to the termination of the evaluation, since higher dimensions would result in exceptionally long run-times.

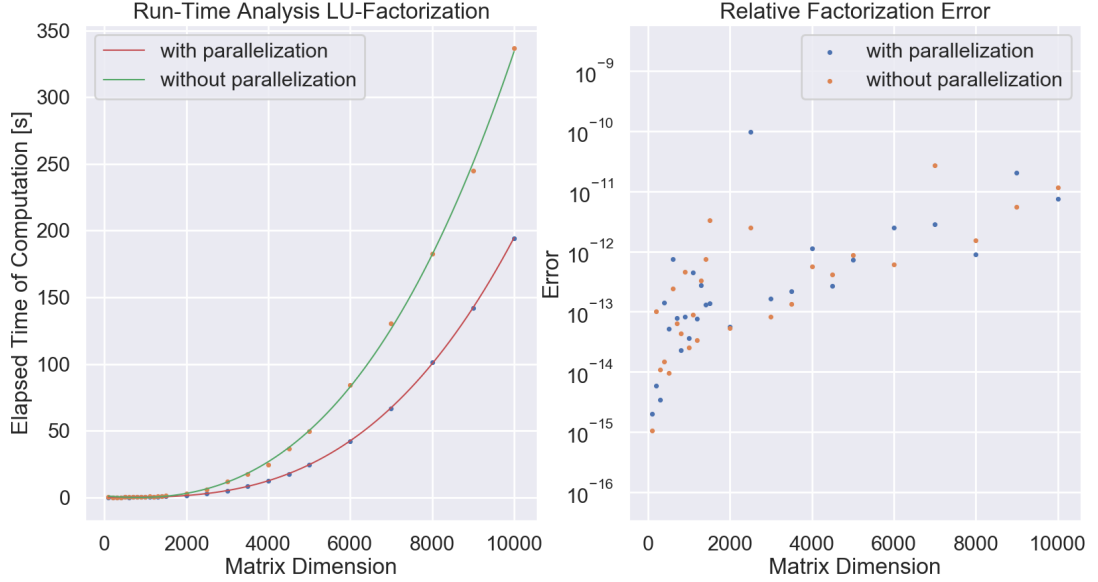


Figure 1: Run-time and Error Analysis

To minimize the run-time of the algorithm, the Python module `< numba >`, which translates the code into much faster machine code, was used. The code was adapted, such that it was possible to use parallelization (via the `parallel=True` argument and the `prange`-function, which indicates *numba* where to apply parallelization) and thus, utilizing all cores of the CPU instead of just one. This resulted in a constant CPU-load of 100% during the computation and increased the computation speed by a factor of approx. 1.5. The comparison between running the code with and without parallelization, is visualized in fig. 1.

Fig. 1 also shows the relative factorization error of the computed solutions w.r.t the matrix dimension. Both, the parallelization and the non-parallelization method, feature comparable factorization errors. At the maximum dimension of 10000, the results still feature relative factorization errors of $\sim 10^{-11}$. Since the accuracy of both methods are approximately equal, the parallelization method is superior, due

to its lower run-time and more efficient use of CPU-capacity.

In order to provide a reference value for the run-time, the code was executed on the server the university provided. Fig. 2 showcases the results. Since the server doesn't feature the module *numba*, the non-parallelization method was evaluated and completed the task for size 5000×5000 in around 340 seconds. The direct comparison of run-times of the non-parallelization method of the server and the user-case (which included the *numba*-module and needed 49 seconds), shows that it completed the task approx. seven times faster. Since two different hardware-setups were used, this value should only be taken as a approximation of the optimization of *numba*, but can be seen as an improvement since other evaluations of identical code, have shown that the user-setup was generally only twice as fast as the server.




Figure 2: Run-time and Error Analysis using the provided university server